

```
In [1]: import pandas as pd
        from sklearn.linear_model import LinearRegression
        from sklearn.metrics import mean_squared_error
        from sklearn import linear_model

data = pd.read_csv("2008_births.csv")
data
```

	INST	RPLACE	RCOUNTY	PLURAL	BDATE	BMONTH	BDAY	BYEAR	SEX	RACE	...	MOTHE
0	1	6800	68	1	2008-01-01	1	1	2008	2	1	...	
1	1	160	1	1	2008-01-02	1	2	2008	2	2	...	
2	1	190	1	1	2008-01-02	1	2	2008	1	1	...	
3	1	4100	41	1	2008-01-03	1	3	2008	2	1	...	
4	1	160	1	1	2008-01-03	1	3	2008	2	1	...	
...	
133417	1	2000	20	1	2008-12-19	12	19	2008	1	1	...	
133418	1	2000	20	1	2008-12-22	12	22	2008	2	1	...	
133419	1	2600	26	1	2008-12-26	12	26	2008	1	1	...	
133420	1	2000	20	1	2008-12-30	12	30	2008	2	1	...	
133421	1	8100	81	1	2008-12-31	12	31	2008	1	0	...	

133422 rows x 125 columns

```
In [2]: import numpy as np
        from sklearn.model_selection import train_test_split
```

```
In [3]: #birthweight is what we want to predict - change this to single target
        birth_weight = data[['BPOUND', 'BOUNCE']]
        birth_weight
```

	BPOUND	BOUNCE
0	4	1
1	8	3
2	9	0
3	7	6
4	9	7
...
133417	6	8
133418	9	2
133419	8	7
133420	5	13
133421	7	3

133422 rows x 2 columns

```
In [4]: #PCA might be a good technique to select predictors

        #note that PCA performs best when data is normalized (range b/w 0 and 1)

        #It is possible to use categorical and continuous predictors
        #for a regression problem. My understanding is you need to make
        #dummy variables for the binary predictors.

        #Variables that we will need to deal with:
        # BDATE, HISPOMOM, HISPDAAD
```

```
In [5]: #Attempting PCA on data
        #for now I drop the BDATE, HISPOMOM AND HISPDAAD
        data_drop = data.drop(["BDATE", "HISPOMOM", "HISPDAAD", "BOUNCE", "BPOUND"], axis = 1)
```

```
In [6]: #get a list of columns in pandas object
        names_of_data = data_drop.columns.tolist()

        #shuffle = false prevents data split being different everytime
        X_train, X_test, y_train, y_test = train_test_split(data_drop, birth_weight, test_size=0.25, shuffle=False)

        #split test into validate and test, again making sure the data is always the same for
        #X_test, X_val, y_test, y_val = train_test_split(X_test, y_test, test_size=0.25, shuffle=False)

        #Normalizing the data
        from sklearn.preprocessing import StandardScaler
        sc = StandardScaler()
        X_train = sc.fit_transform(X_train)
        X_test = sc.transform(X_test)

        #running the actual PCA
        from sklearn.decomposition import PCA

        pca = PCA()
        X_train = pca.fit_transform(X_train)
        X_test = pca.transform(X_test)

        #relief f algorithm - sorting features
```

```
In [7]: explained_variance = pca.explained_variance_ratio_
        print(len(explained_variance))
        print(explained_variance)

120
[3.89856404e-02 3.41089470e-02 3.05310536e-02 2.86690661e-02
 2.39828710e-02 2.07124228e-02 1.81565355e-02 1.70798239e-02
 1.68763289e-02 1.59961073e-02 1.56220702e-02 1.35585642e-02
 1.30251204e-02 1.13087563e-02 1.10457470e-02 1.09427186e-02
 1.06522571e-02 1.03091115e-02 1.02034145e-02 1.01603763e-02
 9.98586972e-03 9.86404012e-03 9.78778404e-03 9.61165628e-03
 9.46902421e-03 9.40630729e-03 9.25092446e-03 9.20797437e-03
 9.16882901e-03 9.11728971e-03 9.08471022e-03 9.05880935e-03
 8.91090184e-03 8.85632587e-03 8.83902308e-03 8.82366452e-03
 8.73228213e-03 8.70972804e-03 8.64475483e-03 8.63888132e-03
 8.60300393e-03 8.57402898e-03 8.54542908e-03 8.51710741e-03
 8.50298288e-03 8.46111398e-03 8.42122923e-03 8.39505222e-03
 8.37548109e-03 8.34046815e-03 8.29732609e-03 8.28947627e-03
 8.24895028e-03 8.22904830e-03 8.20238682e-03 8.12690154e-03
 8.11341630e-03 8.08291392e-03 8.07851589e-03 8.03763212e-03
 8.01473052e-03 7.96613523e-03 7.90999598e-03 7.89944166e-03
 7.83600377e-03 7.82191448e-03 7.78512254e-03 7.75691445e-03
 7.69956508e-03 7.66449230e-03 7.60968558e-03 7.58920895e-03
 7.51965207e-03 7.50072699e-03 7.41546041e-03 7.36057792e-03
 7.17224177e-03 7.11636014e-03 7.01568819e-03 6.99311496e-03
 6.88970752e-03 6.80287045e-03 6.71667348e-03 6.59331242e-03
 6.56011619e-03 6.39097514e-03 6.21038587e-03 6.13263995e-03
 6.02101475e-03 5.88755078e-03 5.62716616e-03 5.49427350e-03
 5.42691648e-03 5.30849077e-03 5.16759622e-03 4.77164460e-03
 4.64430993e-03 4.52817477e-03 4.35785968e-03 4.10170975e-03
 3.95902522e-03 3.70091254e-03 3.10259706e-03 2.89941030e-03
 2.48400831e-03 1.79876640e-03 1.72329185e-03 9.56485073e-04
 9.01290670e-04 1.5880756e-04 5.54770674e-04 2.94096489e-04
 1.90827613e-04 3.83559595e-08 7.68856188e-33 7.62231993e-33
 2.16597016e-33 1.43573650e-33 2.23811207e-34 1.98989560e-34]
```

```
In [8]: #Explained variance prints the variance each principal component contributes.
        #As we can see, the last 5 contribute very little (maybe we can get rid of?)

        #We also want to check for linearity between the input predictors and the output
        #If there is high colinearity, then we want to use ridge regression - A variant of linear regression
        #Correlation indicates strength and direction of a linear relationship. let's use this
```

```
In [9]: # Output of Y into pounds by adding the pounds and ounces as pounds

        y_train = y_train['BPOUND'] + y_train['BOUNCE']*0.0625
        y_test = y_test['BPOUND'] + y_test['BOUNCE']*0.0625
```

This will be multiple linear regression below

```
In [10]: #Setting up the linear regression and fitting the model
        regressFunc = linear_model.LinearRegression(copy_X = True, fit_intercept = True, n_jobs=-1)
        regressFunc.fit(X_train, y_train)
        #Predict the values of the test set
        predicted_val = regressFunc.predict(X_test)
```

```
In [11]: #Comparing the results vs. expected
        y_test,predicted_val
```

Out[11]:	(106737 7.8125 106738 7.6875 106739 5.1250 106740 8.0000 106741 6.3750 ... 133417 6.5000 133418 9.1250 133419 8.4375 133420 5.8125 133421 7.1875 Length: 26685, dtype: float64, array([8.23676507, 7.25996006, 5.0400839 , ..., 6.08695081, 3.87761499, 5.00857935]))
----------	--

```
In [12]: mean_squared_error(y_test, predicted_val)
```

Out[12]: 7.178357456009061e+23

```
In [ ]:
```